

Recitation 10: Malloc Lab

Instructors

Nov. 5, 2018

Administrivia

- **Malloc checkpoint due Thursday, Nov. 8! woooooooooooo**
- **Malloc final due the week after, Nov. 15! woooooooooooo**
- **Malloc Bootcamp Sunday, Nov. 11 at Rashid Auditorium, 7-8:30PM**
 - We will cover fun and flirty ways to succeed post-malloc checkpoint!
 - Tell your friends to come (if they're in 213 (if they want to come (don't force your friends to do things they don't want to do that's not what friends are for)))

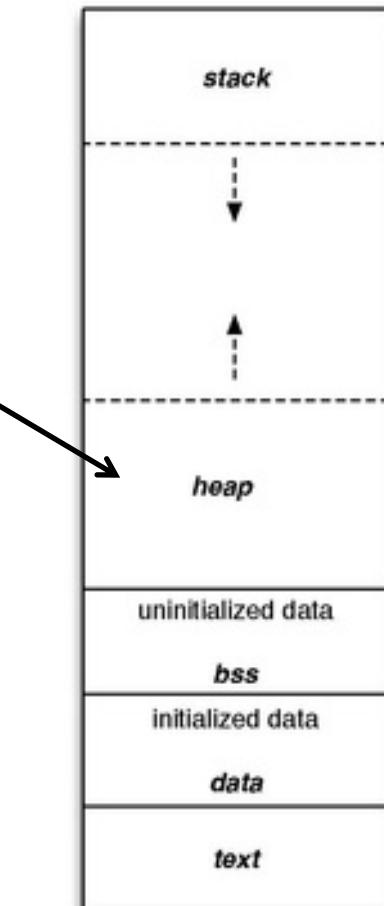
Outline

- Concept
- How to choose blocks
- Metadata
- Debugging / GDB Exercises

What is malloc?

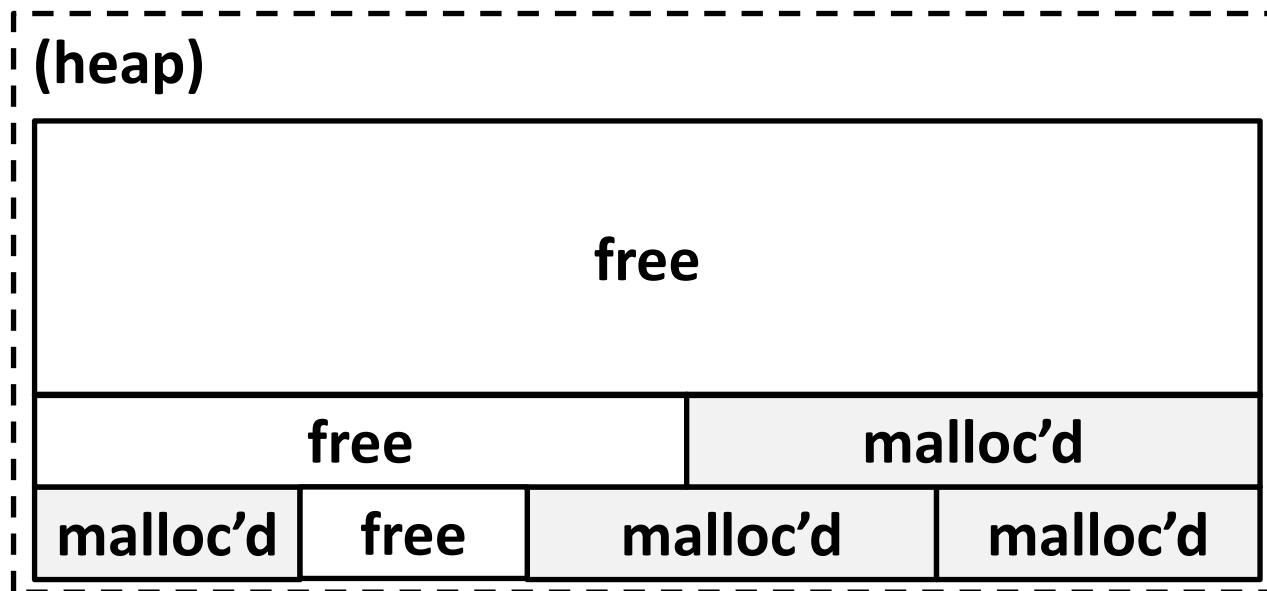
- A function to allocate memory during runtime (dynamic memory allocation).
 - More useful when the size or number of allocations is unknown until runtime (e.g. data structures)

- The heap is a segment of memory addresses reserved almost exclusively for malloc to use.
 - Your code directly manipulates the bytes of memory in this section.



Malloc Internals

- The heap consists of blocks of memory



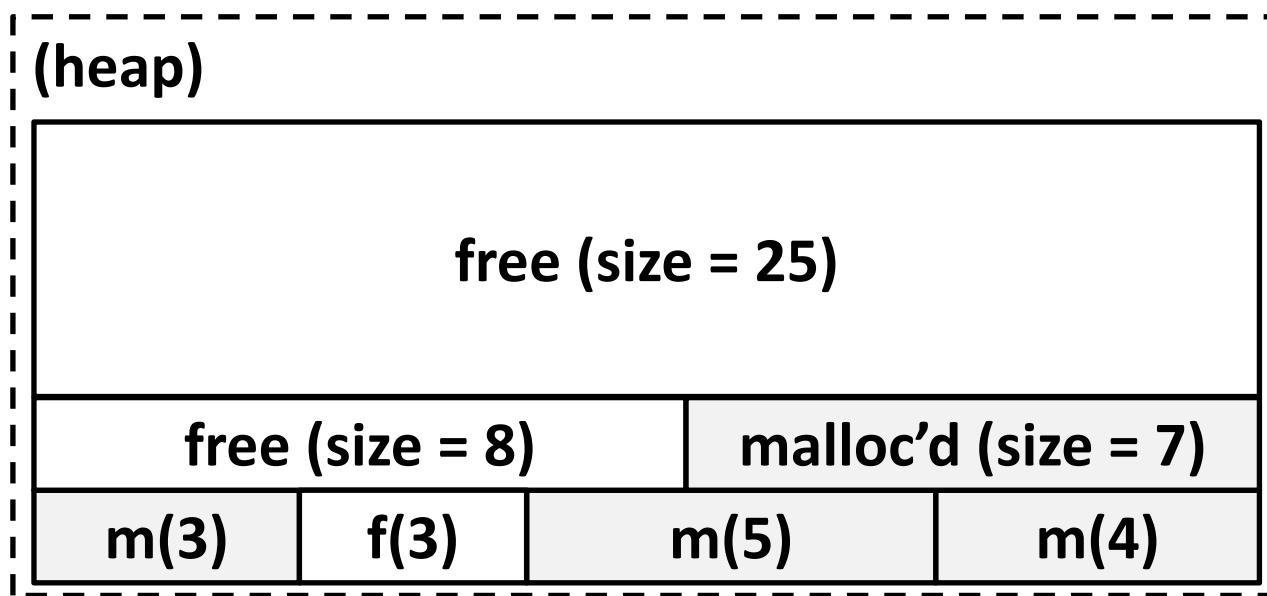
Concept

- Overall, malloc does three things:
 1. Organizes all blocks and stores information about them in a structured way.
 2. Uses the structure made to choose an appropriate location to allocate new memory.
 3. Updates the structure when the user frees a block of memory.

This process occurs even for a complicated algorithm like segregated lists.

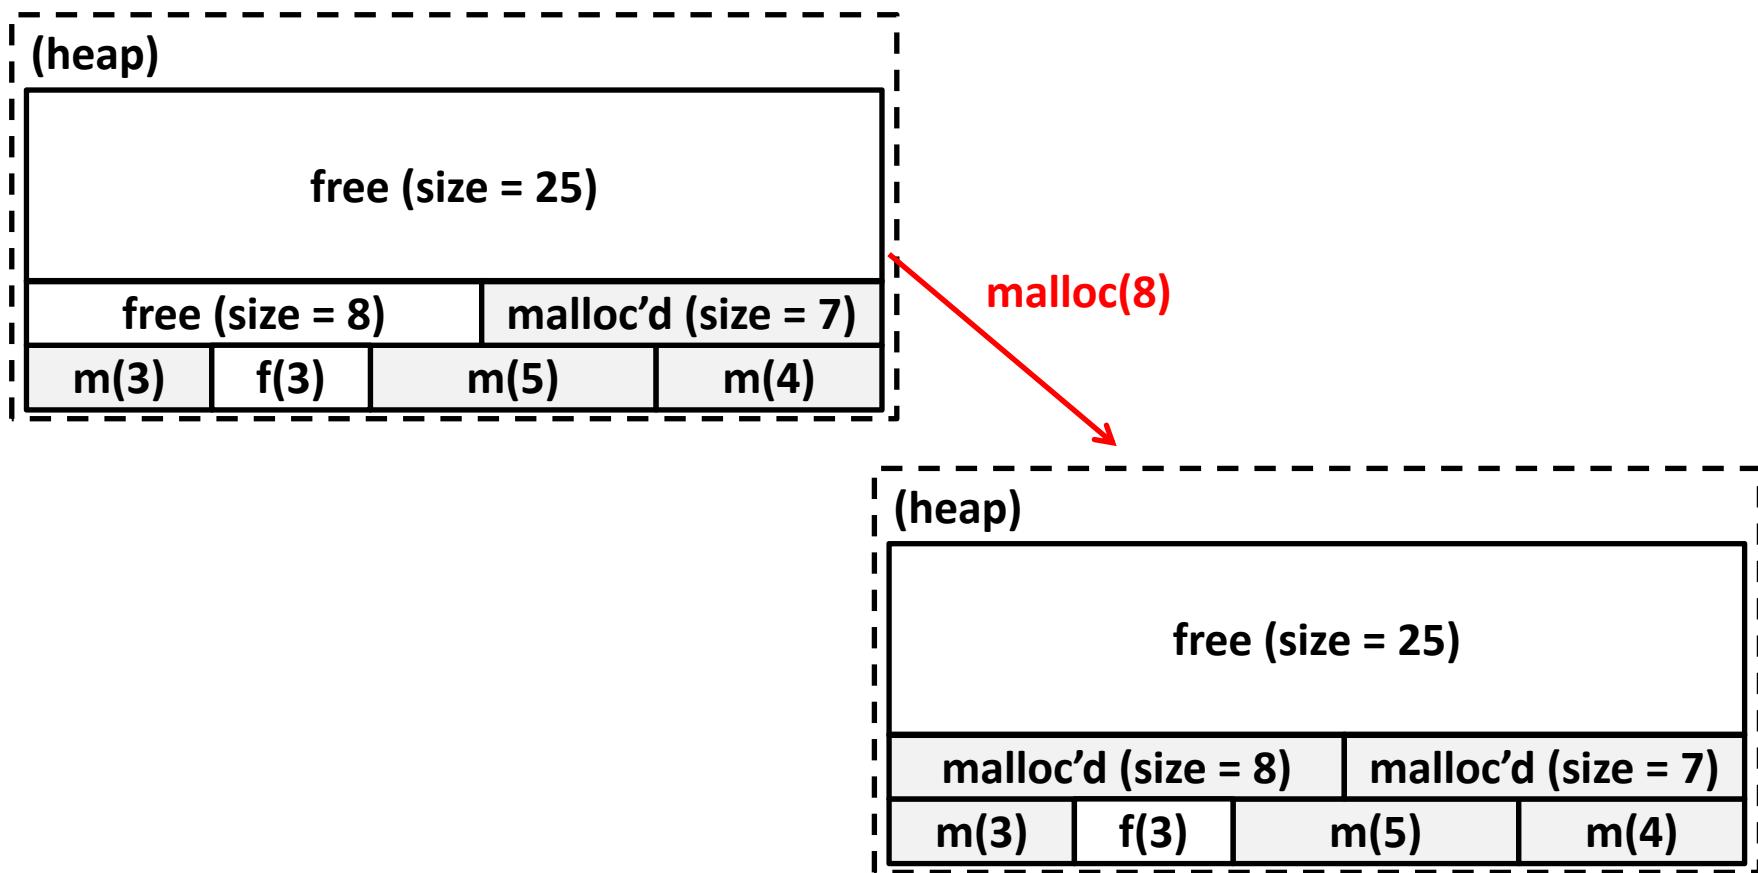
Concept (Implicit list)

1. Organizes all blocks and stores information about them in a structured way.



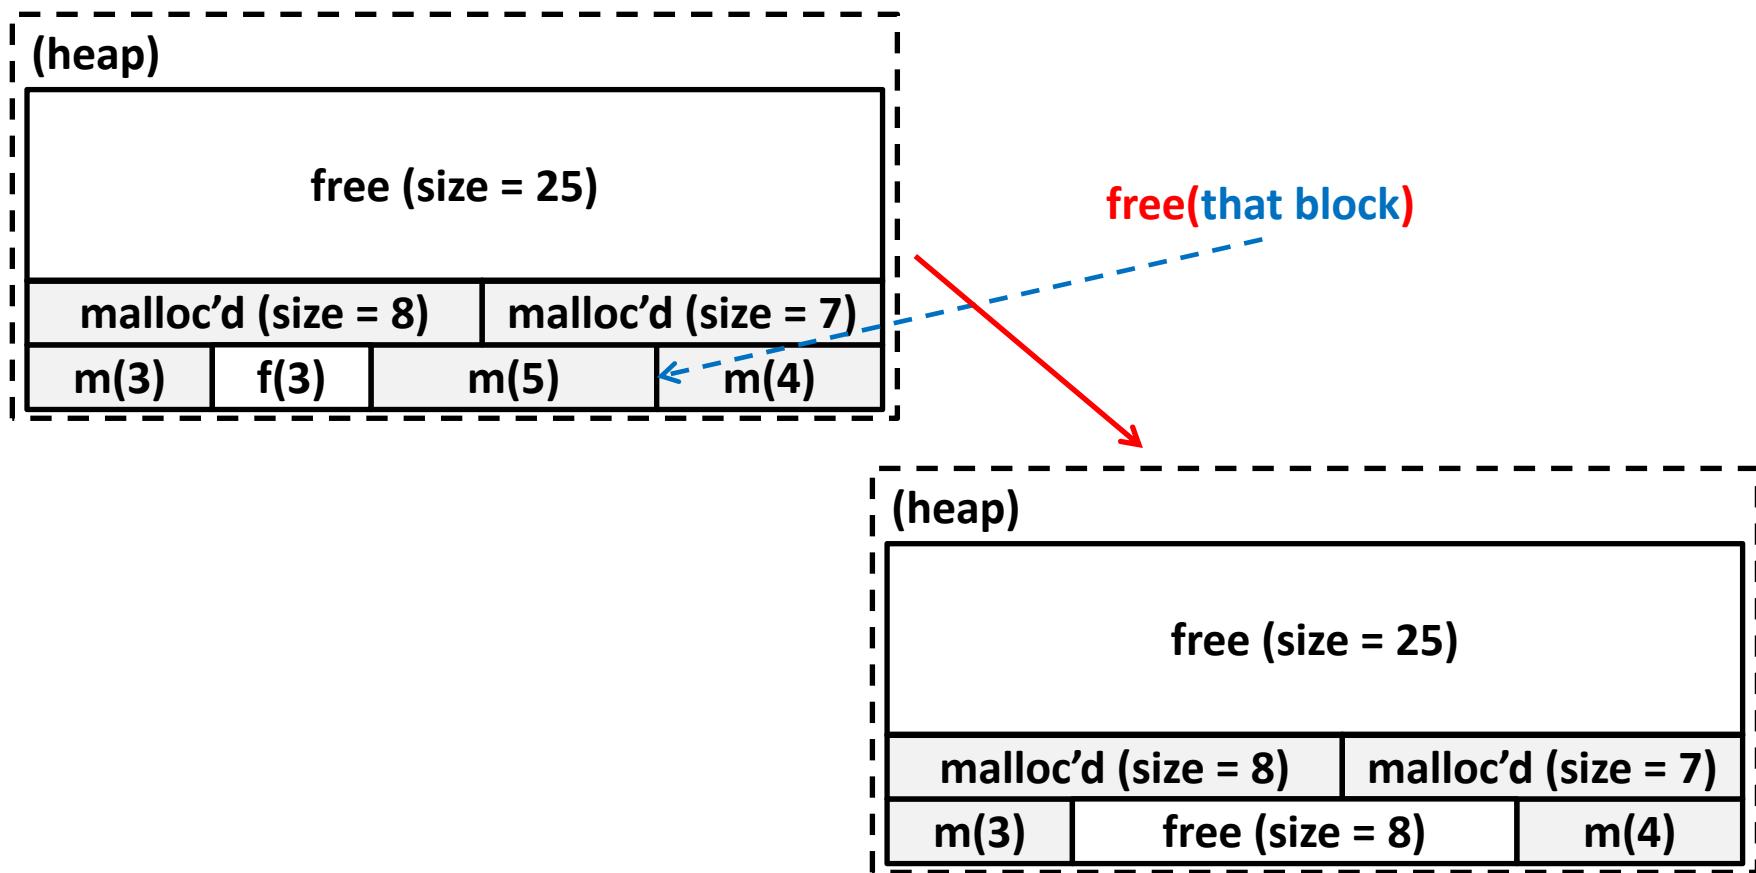
Concept (Implicit list)

2. Uses the structure made to choose an appropriate location to allocate new memory.



Concept (Implicit list)

- Updates the structure when the user frees a block of memory.



Goals

- Run as fast as possible
- Waste as little memory as possible
- Seemingly conflicting goals, but with ~~the library malloc call~~ cleverness you can do very well in both areas!
- The simplest implementation is the implicit list.
mm-baseline uses this method.
 - Unfortunately...

```
[dalud@angelshark:~/.../15213/s17/malloclabcheckpoint-handout] $ ./mdriver -p
Found benchmark throughput 13090 for cpu type Intel(R)Xeon(R)CPUE5520@2.27GHz, benchmark checkpoint
Throughput targets: min=2618, max=11781, benchmark=13090
.....
Results for mm malloc:
  valid   util     ops    msecs   Kops   trace
  yes    78.4%      20     0.002   9632 ./traces/syn-array-short.rep
  yes    13.4%      20     0.001  25777 ./traces/syn-struct-short.rep
  yes    15.2%      20     0.001  24783 ./traces/syn-string-short.rep
  yes    73.1%      20     0.001  19277 ./traces/syn-mix-short.rep
  yes    16.0%      36     0.001  31192 ./traces/ngram-foxl.rep
  yes    73.6%     757     0.145   5237 ./traces/syn-mix-realloc.rep
* yes   62.0%     5748     3.925   1464 ./traces/bdd-aa4.rep
* yes   58.3%   87830   1682.766    52 ./traces/bdd-aa32.rep
* yes   58.0%   41080   410.385   100 ./traces/bdd-ma4.rep
* yes   58.1%  115380   4636.711    25 ./traces/bdd-nq7.rep
* yes   56.6%   20547    26.677   770 ./traces/cbit-abs.rep
* yes   55.8%   95276    675.303   141 ./traces/cbit-parity.rep
* yes   58.0%   89623    611.511   147 ./traces/cbit-satadd.rep
* yes   49.6%   50583    185.382   273 ./traces/cbit-xyz.rep
* yes   40.6%   32540    76.919   423 ./traces/ngram-gulliver1.rep
* yes   42.4%  127912   1284.959   100 ./traces/ngram-gulliver2.rep
* yes   39.4%   67012    338.591   198 ./traces/ngram-moby1.rep
* yes   38.6%   94828    701.305   135 ./traces/ngram-shakel.rep
* yes   90.9%  80000   1455.891    55 ./traces/syn-array.rep
* yes   88.0%  80000    915.167    87 ./traces/syn-mix.rep
* yes   74.3%  80000    914.366    87 ./traces/syn-string.rep
* yes   75.2%  80000    812.748    98 ./traces/syn-struct.rep
16 16   59.1% 1148359  14732.604    78

Average utilization = 59.1%. Average throughput = 78 Kops/sec
Checkpoint Perf index = 20.0 (util) + 0.0 (thru) = 20.0/100
```

This is pretty
slow... most
explicit list
implementations
get above 10000
Kops/sec

Allocation methods in a nutshell

- **Implicit list:** a list is implicitly formed by jumping between blocks, using knowledge about their sizes.



- **Explicit list:** Free blocks explicitly point to other blocks, like in a linked list.

- Understanding explicit lists requires understanding implicit lists



- **Segregated list:** Multiple linked lists, each containing blocks in a certain range of sizes.

- Understanding segregated lists requires understanding explicit lists



Choices

■ What kind of implementation to use?

- Implicit list, explicit list, segregated lists, binary tree methods, etc.
- You can use specialized strategies depending on the size of allocations
- Adaptive algorithms are fine, though not necessary to get 100%.
 - Don't directly test for which trace file is running.

■ What fit algorithm to use?

- Best fit: choose the smallest block that is big enough to fit the requested allocation size
- First fit / next fit: search linearly starting from some location, and pick the first block that fits.
- Which is faster? Which uses less memory?
- “Good enough” fit: a blend between the two

■ This lab has many more ways to get an A+ than, say, Cache Lab Part 2

Finding a Best Block

- Suppose you have implemented the explicit list approach
 - You were using best fit with explicit lists
- You experiment with using segregated lists instead.
Still using best fits.
 - Will your memory utilization score improve?

Note: you don't have to implement seglists and run mdriver to answer this. That's, uh, hard to do within one recitation session.

- What other advantages does segregated lists provide?
- Losing memory because of the way you choose your free blocks is called external fragmentation.

Metadata

- All blocks need to store some data about themselves in order for malloc to keep track of them (e.g. headers)
 - This takes memory too...
 - Losing memory for this reason is called internal fragmentation.
- What data might a block need?
 - Does it depend on the malloc implementation you use?
 - Is it different between free and allocated blocks?
- Can we use the extra space in free blocks?
 - Or do we have to leave the space alone?
- How can we overlap two different types of data at the same location?

In a perfect world...

Setting up the blocks, metadata, lists... etc (500 LoC)

- + Finding and allocating the right blocks (500 LoC)**
- + Updating your heap structure when you free (500 LoC) =**

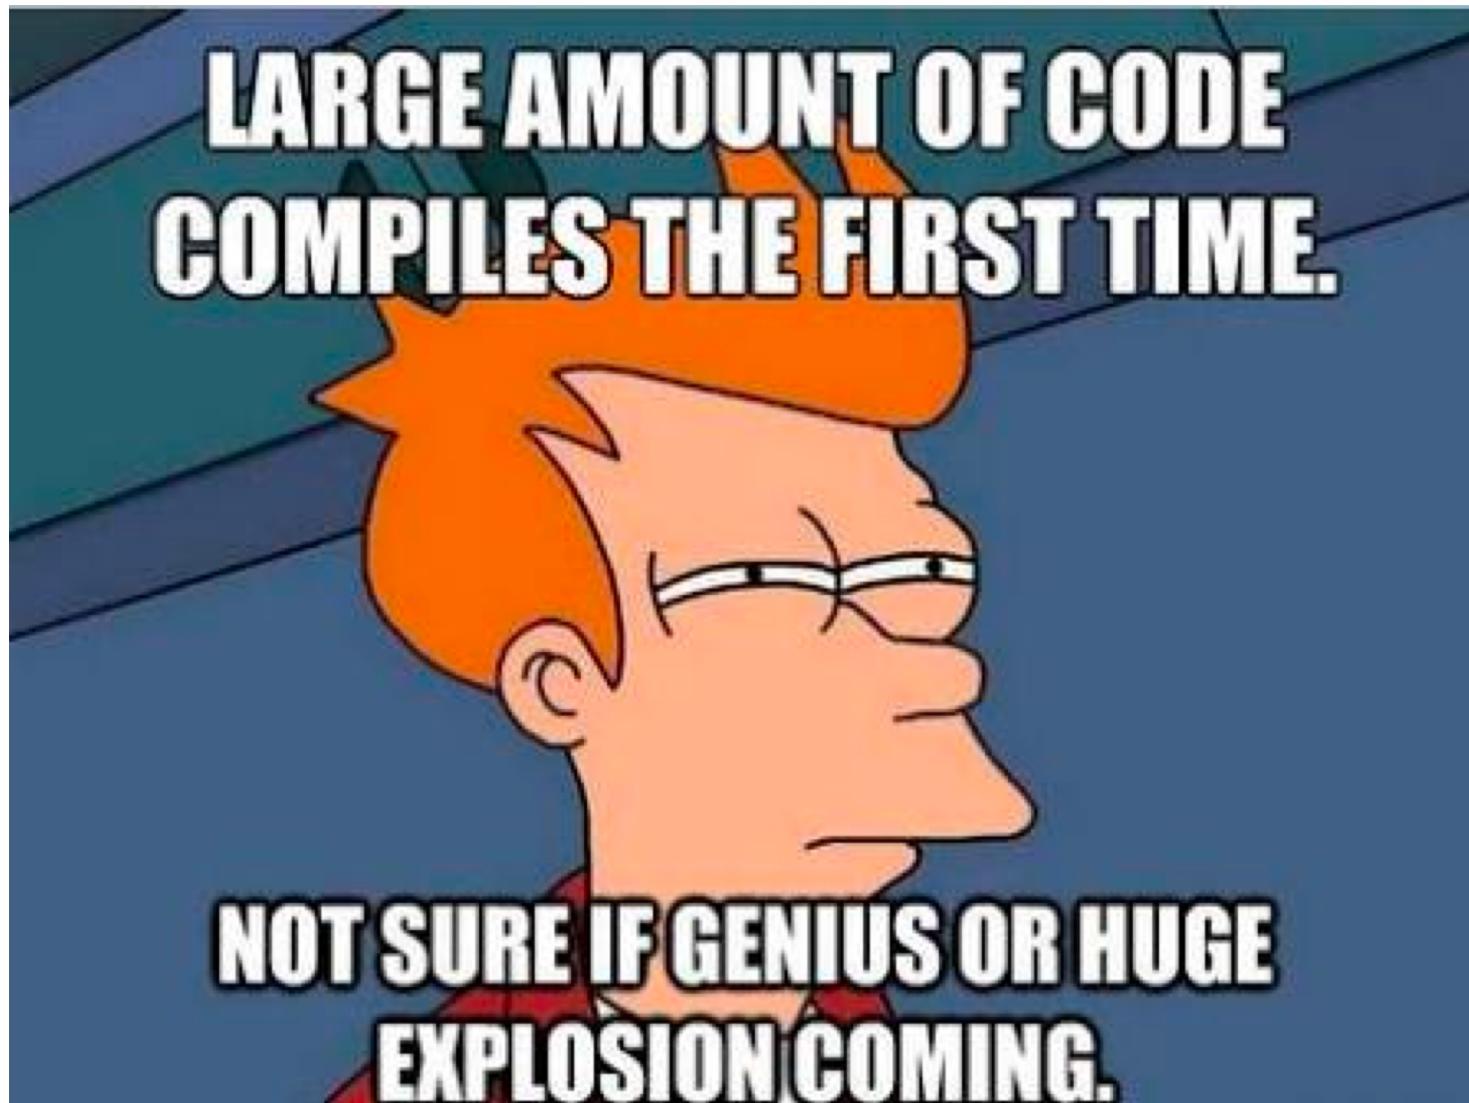
```
[dalud@angelshark:~/.../15213/s17/malloclabcheckpoint-handout] $ ./mdriver
Found benchmark throughput 13056 for cpu type Intel(R) Xeon(R) CPU E5520 @ 2.27G
Throughput targets: min=6528, max=11750, benchmark=13056
.....
Results for mm malloc:
  valid    util      ops     msecs     Kops   trace
    yes    78.1%      20      0.004    5595  ./traces/syn-array-short.rep
    yes    3.2%       20      0.004    5273  ./traces/syn-struct-short.rep
  * yes   96.0%    80000     17.176    4658  ./traces/syn-array.rep
  * yes   93.2%    80000      6.154   12999  ./traces/syn-mix.rep
  * yes   86.4%    80000      3.717   21521  ./traces/syn-string.rep
  * yes   85.6%    80000      3.649   21924  ./traces/syn-struct.rep
16 16    74.2% 1148359     55.949   20525

Average utilization = 74.2%. Average throughput = 20525 Kops/sec
Perf index = 60.0 (util) + 40.0 (thru) = 100.0/100
```

In reality...

- Setting up the blocks, metadata, lists... etc (500 LoC)
- + Finding and allocating the right blocks (500 LoC)
- + Updating your heap structure when you free (500 LoC)
- + One bug, somewhere lost in those 1500 LoC =

```
[dalud@angelshark:~/.../15213/s17/malloclabcheckpoint-handout] $ ./mdriver
Found benchmark throughput 13056 for cpu type Intel(R) Xeon(R) CPU E5520 @ 2.27GHz
Throughput targets: min=6528, max=11750, benchmark=13056
.....Segmentation fault
[dalud@angelshark:~/.../15213/s17/malloclabcheckpoint-handout] $ █
```



Common errors you might see

■ Garbled bytes

- Problem: overwriting data in an allocated block
- Solution: ~~remembering data lab and the good ol' days~~ finding where you're overwriting by stepping through with gdb

■ Overlapping payloads

- Problem: having unique blocks whose payloads overlap in memory
- Solution: ~~literally print debugging everywhere~~ finding where you're overlapping by stepping through with gdb

■ Segmentation fault

- Problem: accessing invalid memory
- Solution: ~~crying a little~~ finding where you're accessing invalid memory by stepping through with gdb

GDB Practice

- Using GDB well in malloclab can save you ***HOURS^{1, 2}*** of debugging time

- Average 20 hours using GDB for “B” on malloclab
 - Average 23 hours not using GDB for “B” on malloclab

- Form pairs

- Login to a shark machine
 - wget <http://www.cs.cmu.edu/~213/activities/rec11.tar>
 - tar xf rec11.tar
 - cd rec11
 - make

- Two buggy mdrivers

1. Average time is based on Summer 2016 survey results
2. As the TA making these slides, I realize that there's really no way for me to confirm that these stats are true, so by the power of anecdotal evidence, let me tell you about how I was too stubborn to use GDB and didn't even finish malloc while another TA Niko mastered GDB in two weeks and cruised to a 100%. Use GDB!!!

First things first

■ Try running `$ make`

- If you look closely, our code compiles your `malloc` implementation with the `-O3` flag.
- This is an optimization flag. `-O3` makes your code run as efficiently as the compiler can manage, but also makes it horrible for debugging (almost everything is “optimized out”).

```
[dalud@angelshark:~/.../15213/s17/rec11] $ make
gcc -Wall -Wextra -Werror -O3 -g -DDRIVER -Wno-unused-function -Wno-u
./macro-check.pl -f mm.c
clang -Wall -Wextra -Werror -O3 -g -DDRIVER -Wno-unused-function -Wno-u
gcc -Wall -Wextra -Werror -O3 -g -DDRIVER -Wno-unused-function -Wno-u
(gdb) print block
$3 = <optimized out>
(gdb) print asize
$4 = <optimized out>
```

Debugging mdriver

```
$ gdb --args ./mdriver -c traces/syn-mix-short.rep
```

```
(gdb) run
```

```
(gdb) backtrace
```

```
(gdb) list
```

Optional: Type Ctrl-X Ctrl-A to see the source code. Don't linger there for long, since this visual mode is buggy. Type that key combination again to go back to console mode.

- 1) What function is listed on the top of backtrace?**
- 2) What line of code crashed?**
- 3) How did that line cause the crash?**

Debugging mdriver

- **(gdb) x /10gx block**
 - Shows the memory contents within the block
 - In particular, look for the header.
- **Remember the output from (gdb) bt?**
- **(gdb) frame 1**
 - Jumps to the function one level down the call stack (aka the function that called `write_footer`)
 - Ctrl-X, Ctrl-A again if you want to see visuals
- **What was the caller function? What is its purpose?**
 - Was it writing to `block` or `block_next` when it crashed?

Thought process while debugging

- **write_footer crashed because it got the wrong address for the footer...**
- **The address was wrong because the header of the block was some garbage value**
 - Since `write_footer` uses `get_size(block)` after all
- **But why in the world does the header contain garbage??**
 - The crash happened in `place`, which basically splits a free block into two and uses the first one to store things.
 - Hm, `block_next` would be the new block created after the split?
The one on the right?
 - The header would be in the middle of the original free block actually. Wait, but I wrote a new header before I wrote the footer!
 - Right? ...Oh, I didn't. Darn.

Heap consistency checker

- mm-2.c activates debug mode, and so mm_checkheap runs at the beginning and end of many of its functions.

```
106 /*  
107  * If DEBUG is defined, enable printing on dbg_printf and contracts.  
108  * Debugging macros, with names beginning "dbg_" are allowed.  
109  * You may not define any other macros having arguments.  
110 */  
111 #define DEBUG // uncomment this line to enable debugging  
112  
113 #ifdef DEBUG  
114 /* When debugging is enabled, these form aliases to useful functions */  
115 #define dbg_printf( ) printf( VA_ARGS )
```

- The next bug will be a total nightmare to find without this heap consistency checker*.

Now you try debugging this

```
$ gdb --args ./mdriver-2 -c traces/syn-array-short.rep
```

mm_checkheap will fail. What reason does it cite?

Where's the footer? Use x /gx and some arithmetic

Track changes in the header and the footer:

(gdb) watch *[header address]

(gdb) watch *[footer address]

When does the footer's value turn inconsistent? What function was running at the time? Which part of that function was wrong? Use backtrace on each frame.

MallocLab Checkpoint

- Due this Thursday!
- Checkpoint should take a bit less than half of the time
- Read the write-up. Slowly. Carefully.
- Use GDB
- Ask us for debugging help
 - Only after you implement mm_checkheap though

Appendix: Advanced GDB Usage

- **backtrace**: Shows the call stack
- **frame**: Lets you go to one of the levels in the call stack
- **list**: Shows source code
- **print <expression>**:
 - Runs any valid C command, even something with side effects like mm_malloc(10) or mm_checkheap(1337)
- **watch <expression>**:
 - Breaks when the value of the expression changes
- **break <function / line> if <expression>**:
 - Only stops execution when the expression holds true
- **Ctrl-X Ctrl-A** for visualization

Appendix: Building O0

- Edit the file named **Makefile** and make it use -O0

```
4 # Regular compiler
5 CC = gcc
6 # Compiler for mm.c
7 CLANG = clang
8 # Change this to -O0 (big-Oh, numeral zero) if you need to use a debugger on your code
9 COPT = -O0
10 CFLAGS = -Wall -Wextra -Werror $(COPT) -g -DDRIVER -Wno-unused-function -Wno-unused-parameter
11 LIBS = -lm -lrt
12
```

- Then run **\$ make -B**
 - Alternative: **\$ make clean \$ make**
 - Just running **make** won't work because it'll say nothing new needs to be compiled. So we force it to recompile.
- Remember to set it back to -O3 when you're done to test throughput, since -O0 makes your code much slower.

Appendix: hprobe()

- Function to view what's on the heap
- Can be used to look at a block's leading structures, can also be used for the emulated addresses